

CS 149 Operating Systems

Processes


Instructor: Kong Li

Content

- Process Concept
- Context switch
- Process Scheduling
- Operations on Processes
 - fork, exec, wait, exit
- Interprocess Communication (IPC)
 - Shared memory, msg passing, socket, pipe

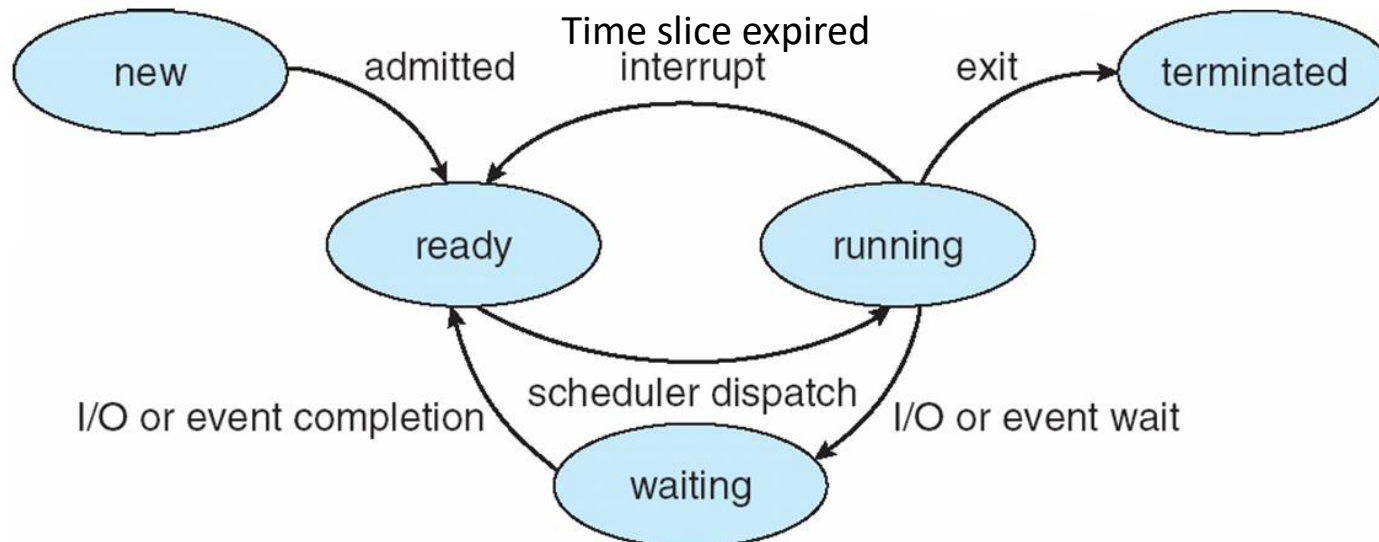
Process Concept

- Program: **passive** entity stored on disk (**executable file**)
 - Program becomes proc when executable file loaded into memory
 - One single program can become several concurrent procs
- Process: **active** entity in memory, a program in execution
 - proc exec: in sequential fashion
- Process parts
 - **program counter**, CPU registers, proc state (later)
 - **Text**: program code
 - **Data**: global variables
 - **Heap**: runtime memory allocation (malloc/free)
 - **Stack**: temporary data
 - Function parameters, return addr, local variables



Process State

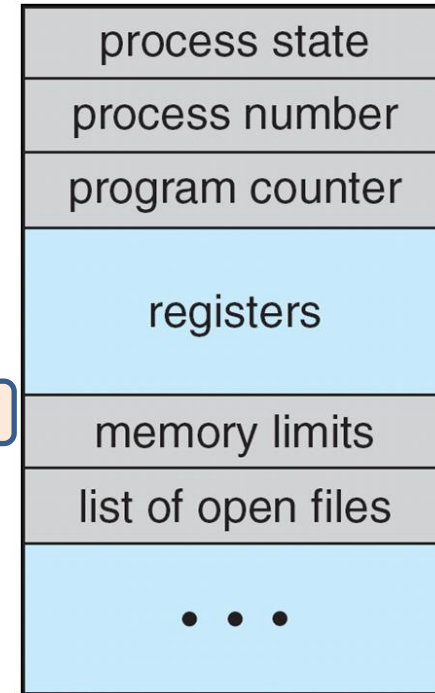
- As a process executes, it changes **state**
 - new: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a processor
 - terminated: The process has finished execution



State transition diagram

Process Control Block (PCB)

- **PCB**: one per proc (aka task control block)
 - Proc state: running, waiting, etc
 - PID: **unique** among all procs on a given computer
 - **Program counter**: next instruction to execute
 - Registers: contents of CPU registers include SP
 - scheduling: priorities, queue pointers
 - Memory-mgmt: memory allocated to proc
 - Accounting: CPU used, clock time elapsed since start, time limits
 - I/O status: – I/O devices allocated to proc, **open files**



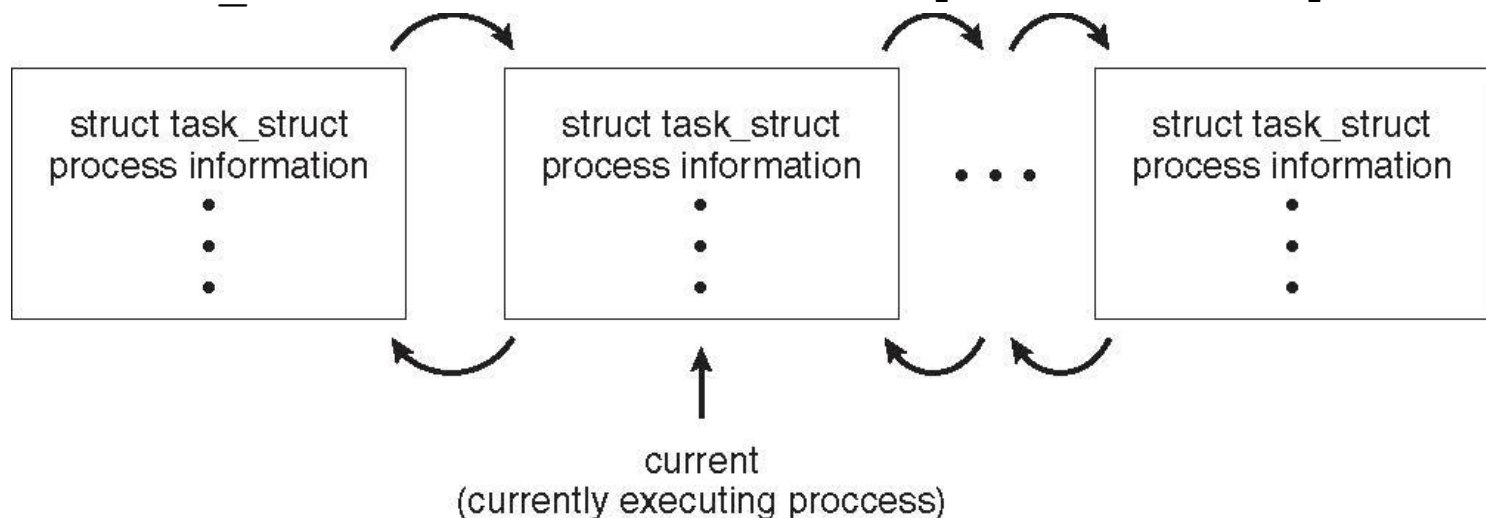
Q: Where is PCB?
A: kernel memory
(in main memory)

Process Representation in Linux

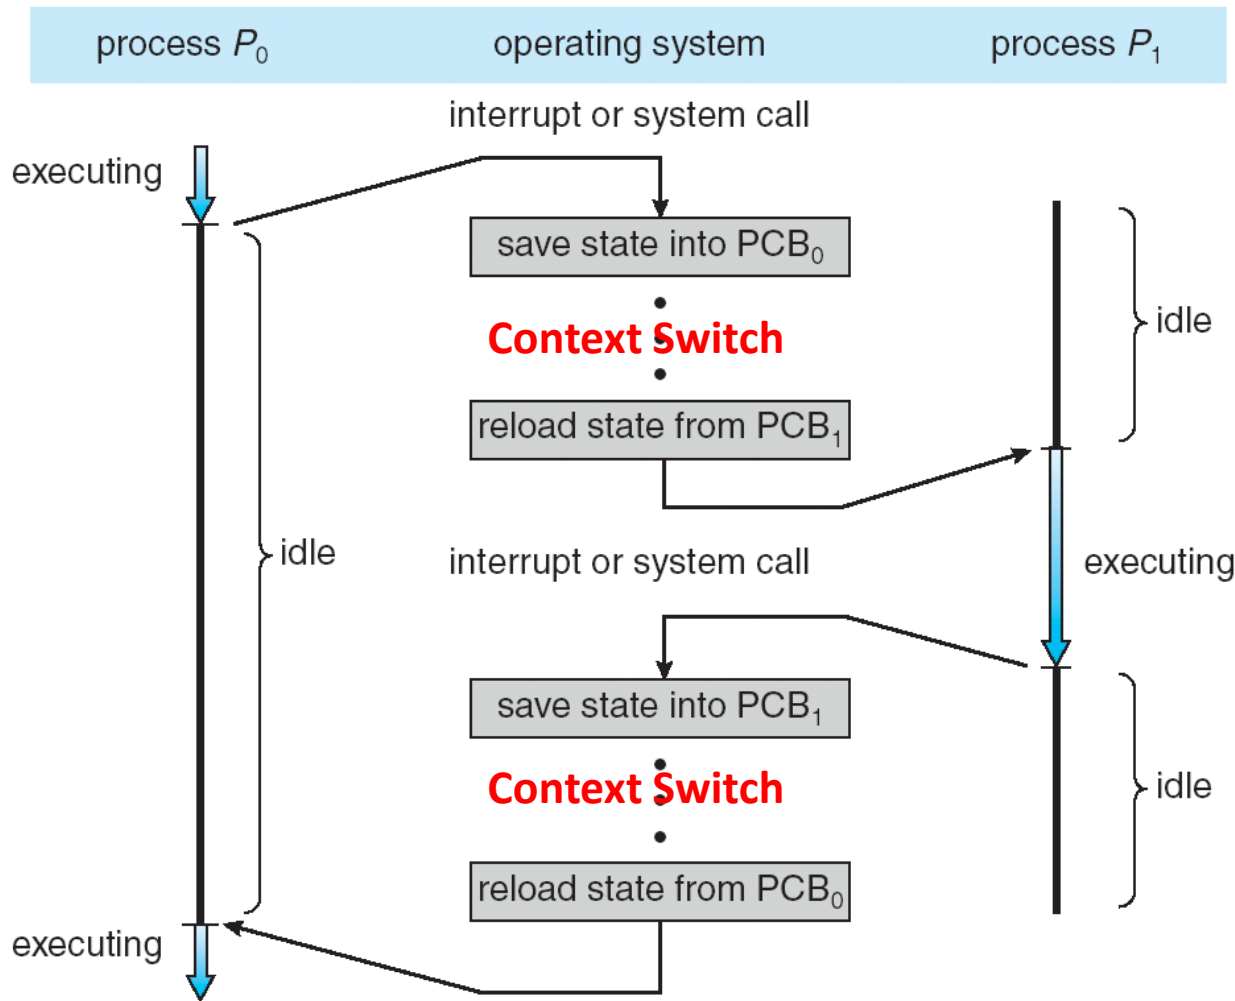
- Represented by the C structure `task_struct`

```
pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this proc */
```

Why?



CPU Switch From Process to Process




Context Switch

- **Context** of a proc represented in the PCB
- **context switch**: CPU switches from one to another proc
 - OS must **save** the state of the **old** proc
 - Then OS must **load** (**restore**) the saved state of the **new** proc
- Context-switch time is **overhead**
 - the system does **no** useful work while switching
 - complicated OS/PCB → longer context switch
 - HW support: multiple sets of registers? Etc.
- How to improve context switch efficiency
 - Reduce the frequency of context switch
 - Speed up context switch: multiple set of HW registers, etc.

Process Scheduling

- Goals: maximize CPU use, quickly switch procs onto CPU for time sharing
- Various scheduling **queues**:
 - Ready Q**: procs that are ready and waiting to execute
 - Device Qs**: set of procs waiting for I/O devices (one Q per device) ←
- Procs migrate among various Qs, based on proc state transition
 - continues this until terminates** (removed from all Qs & deallocated resources)

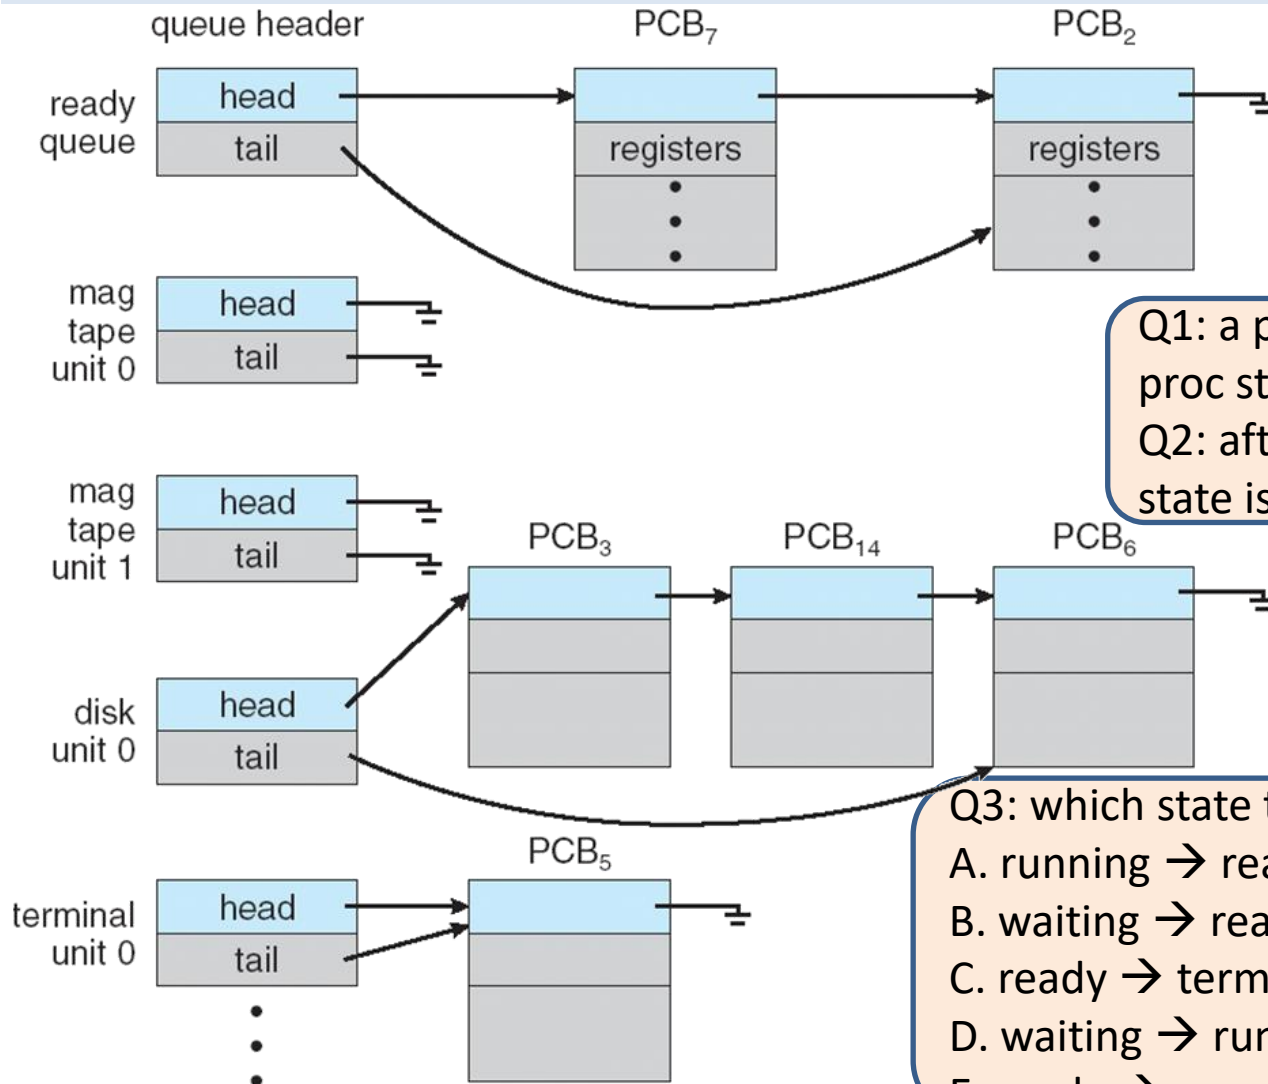


CPU scheduler: runs frequently
Job scheduler: runs less frequently

- Select a good mix of CPU-bound and I/O-bound procs
- Not exist on UNIX & Windows

- CPU (short-term) scheduler**: selects among ready procs for execution
- Job (long-term) scheduler: controls degree of **multiprogramming**

Ready Queue, I/O Device Queues

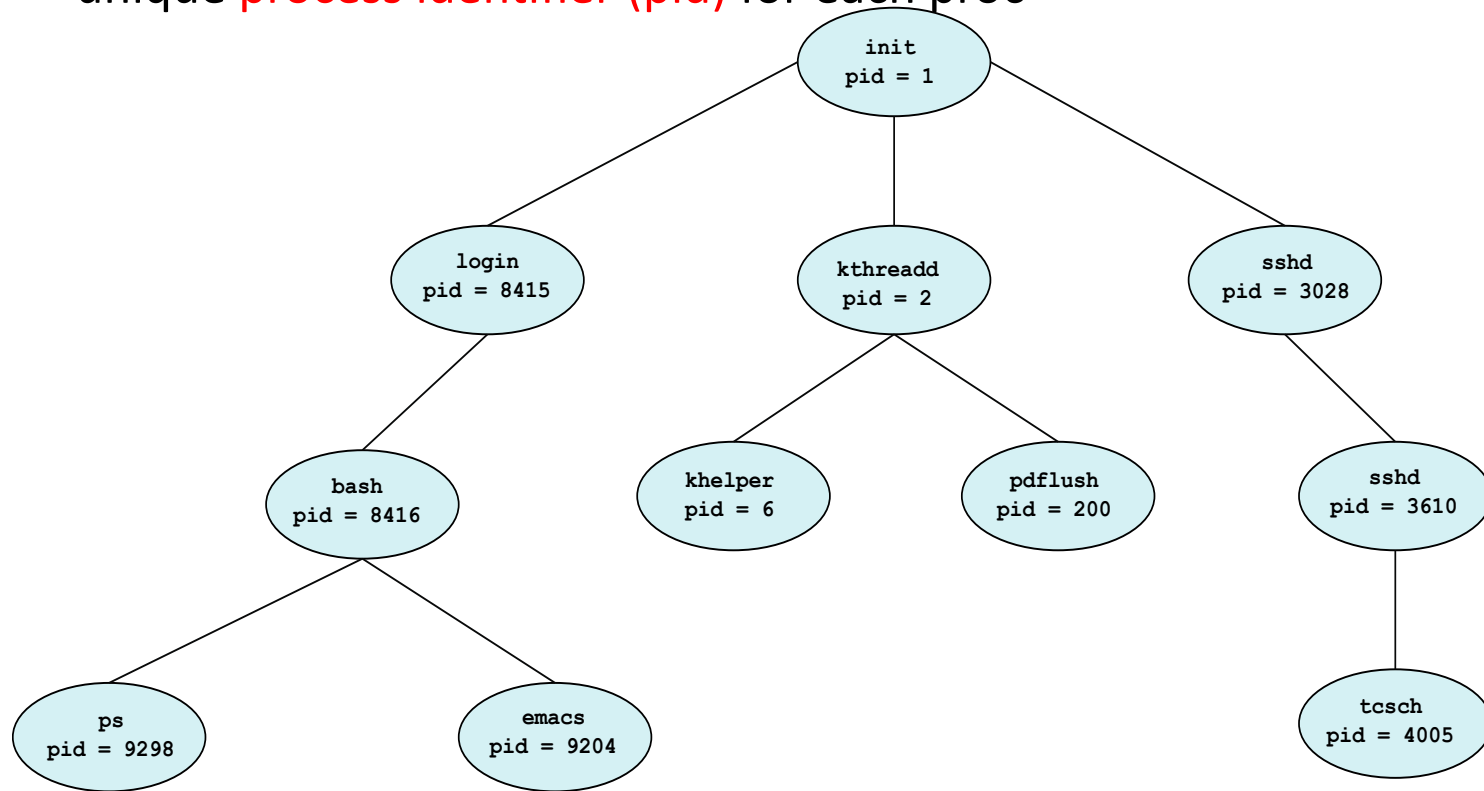


Q1: a proc invokes sleep(10), the proc state is changed from ? to ?
Q2: after 10 seconds, the proc state is changed from ? to ?

Q3: which state transition are not valid?
A. running → ready
B. waiting → ready
C. ready → terminated
D. waiting → running
E. ready → running

Process Tree/Hierarchy

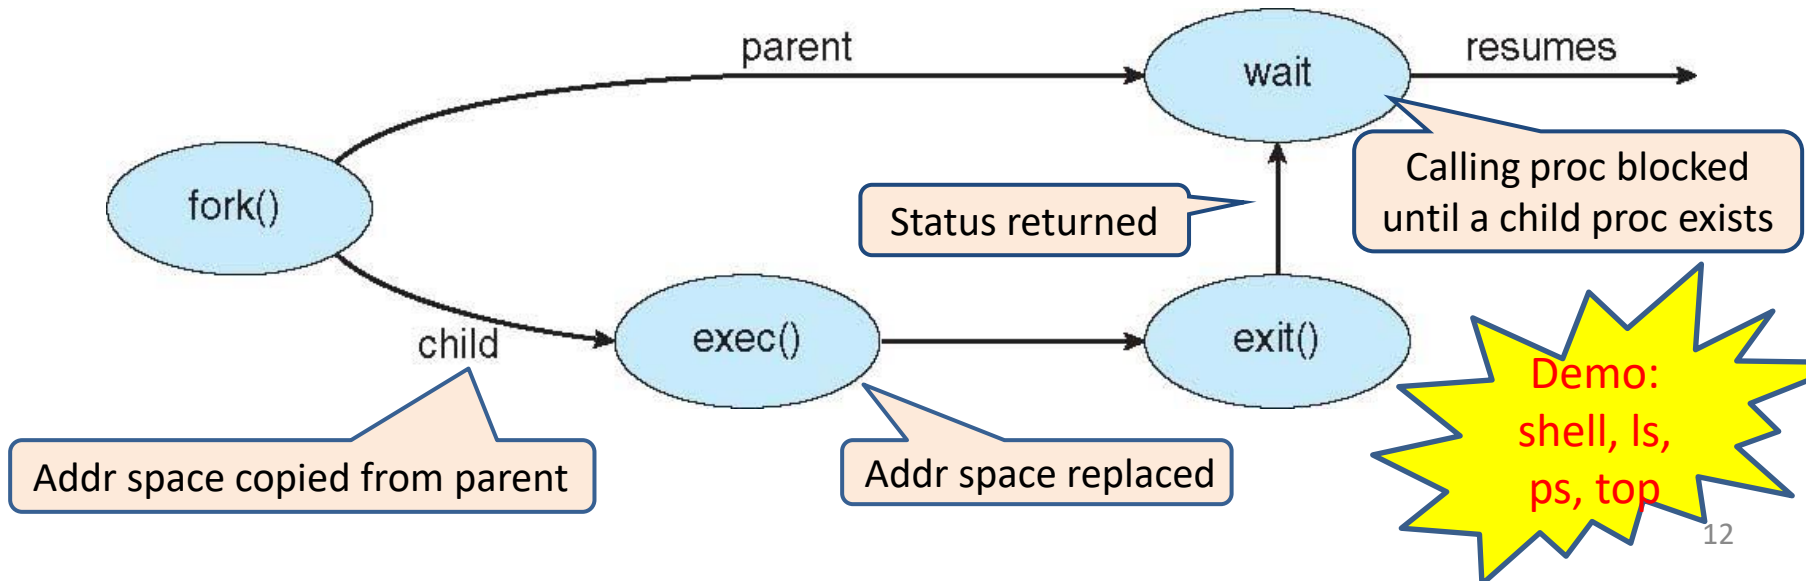
- Linux/UNIX **process tree**
 - **Parent** proc creates **children** procs
 - Child proc creates grandchild procs, etc.
 - PCB records parent proc and children procs
 - unique **process identifier (pid)** for each proc



Process Creation

- UNIX system calls
 - **fork()** : creates new proc – **dup** of parent proc
 - **Both** parent and child procs are ready/running
 - **exec()** : **replace** the proc's memory space w/ a **new** program
 - Usually called after a `fork()`
 - Ex: shell (bash, csh, etc)
- Unlike Windows, no UNIX sys call creates a different new proc

man fork
man exec



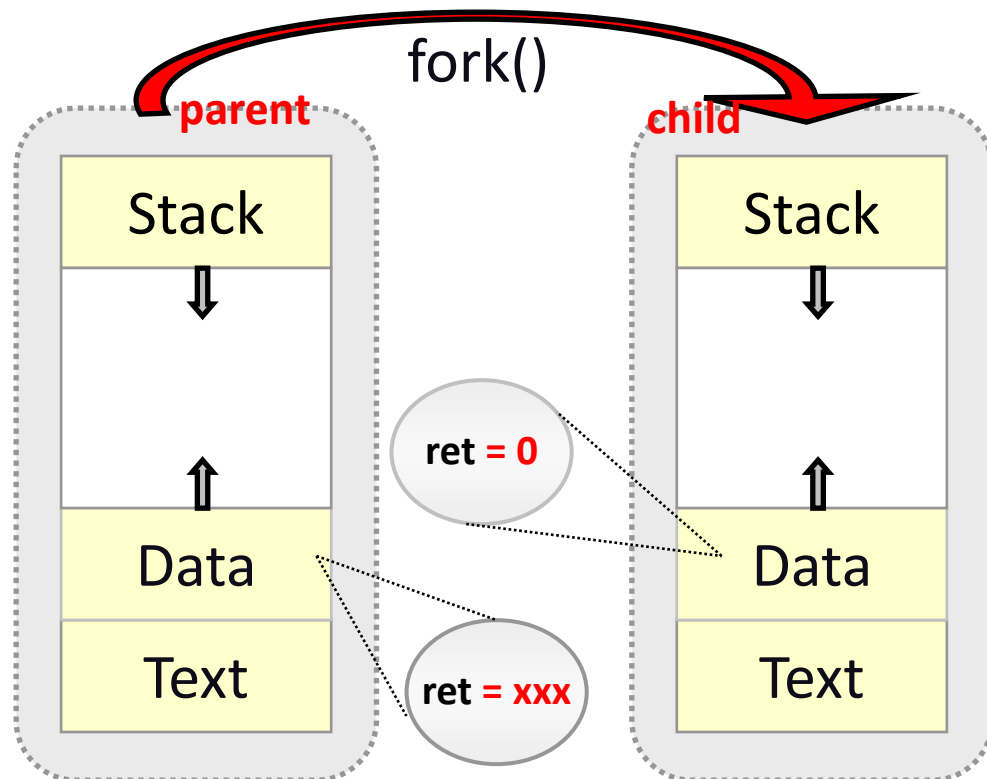
fork()

```
#include <unistd.h>
pid_t fork(void);
```

man fork

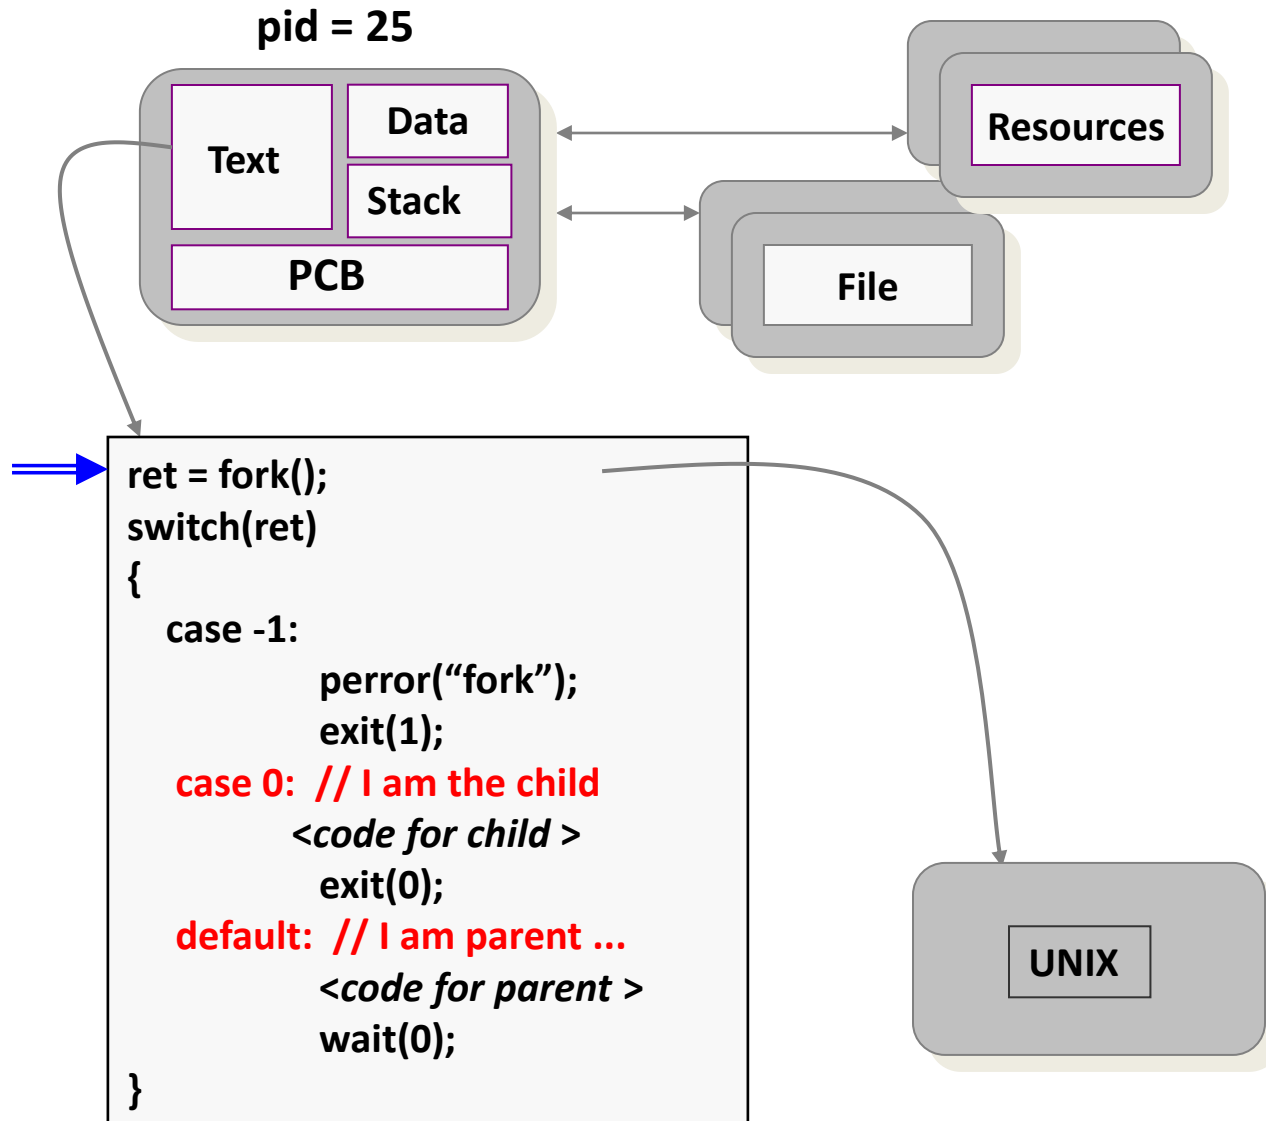
- Current parent process is cloned to a new child proc
 - Both parent and child procs are ready/running
 - Separate memory space
 - The child proc **inherits** memory (text, data, heap, stack) & “open files” from parent (**except?**)
 - **Both procs return from fork()**

```
ret = fork();
ret == -1 if unsuccessful
ret == 0 in the child
ret == child's PID in the parent
```

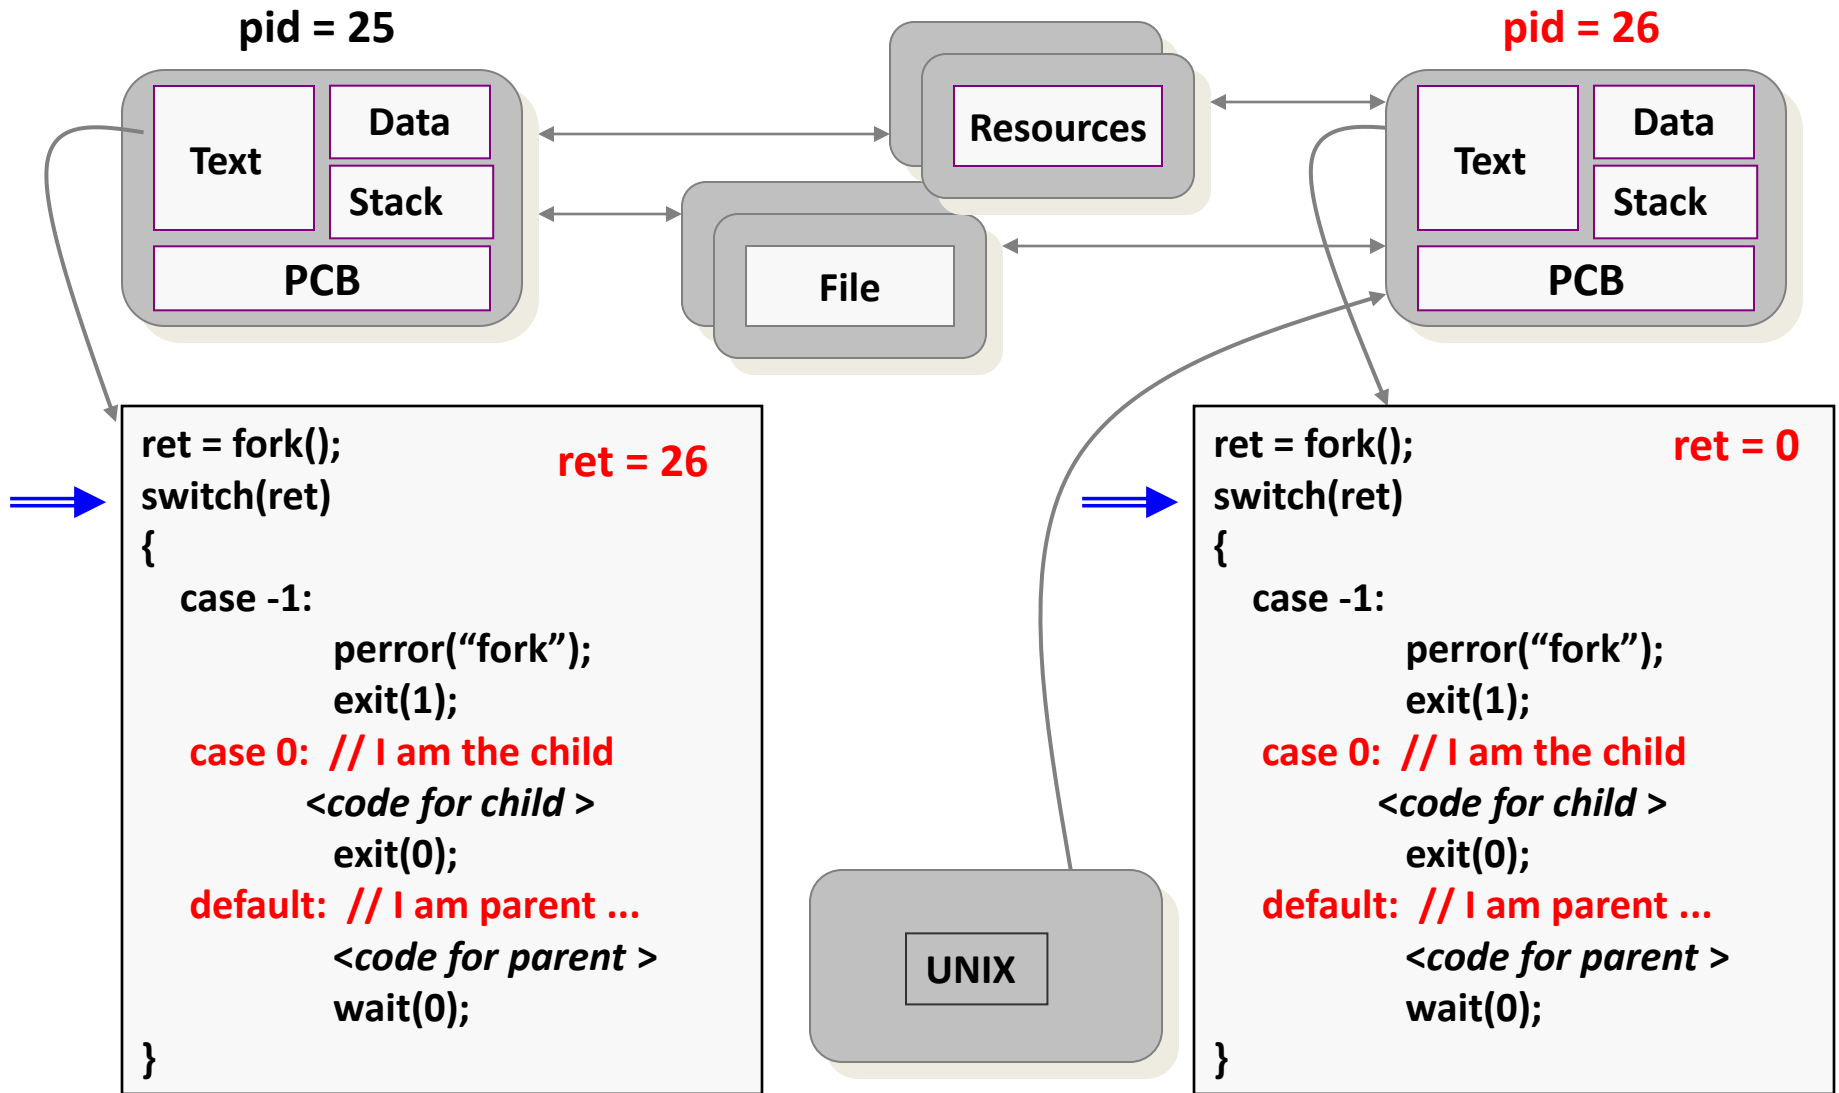


- adapted from Prof. Thomas Way thomas.way@villanova.edu

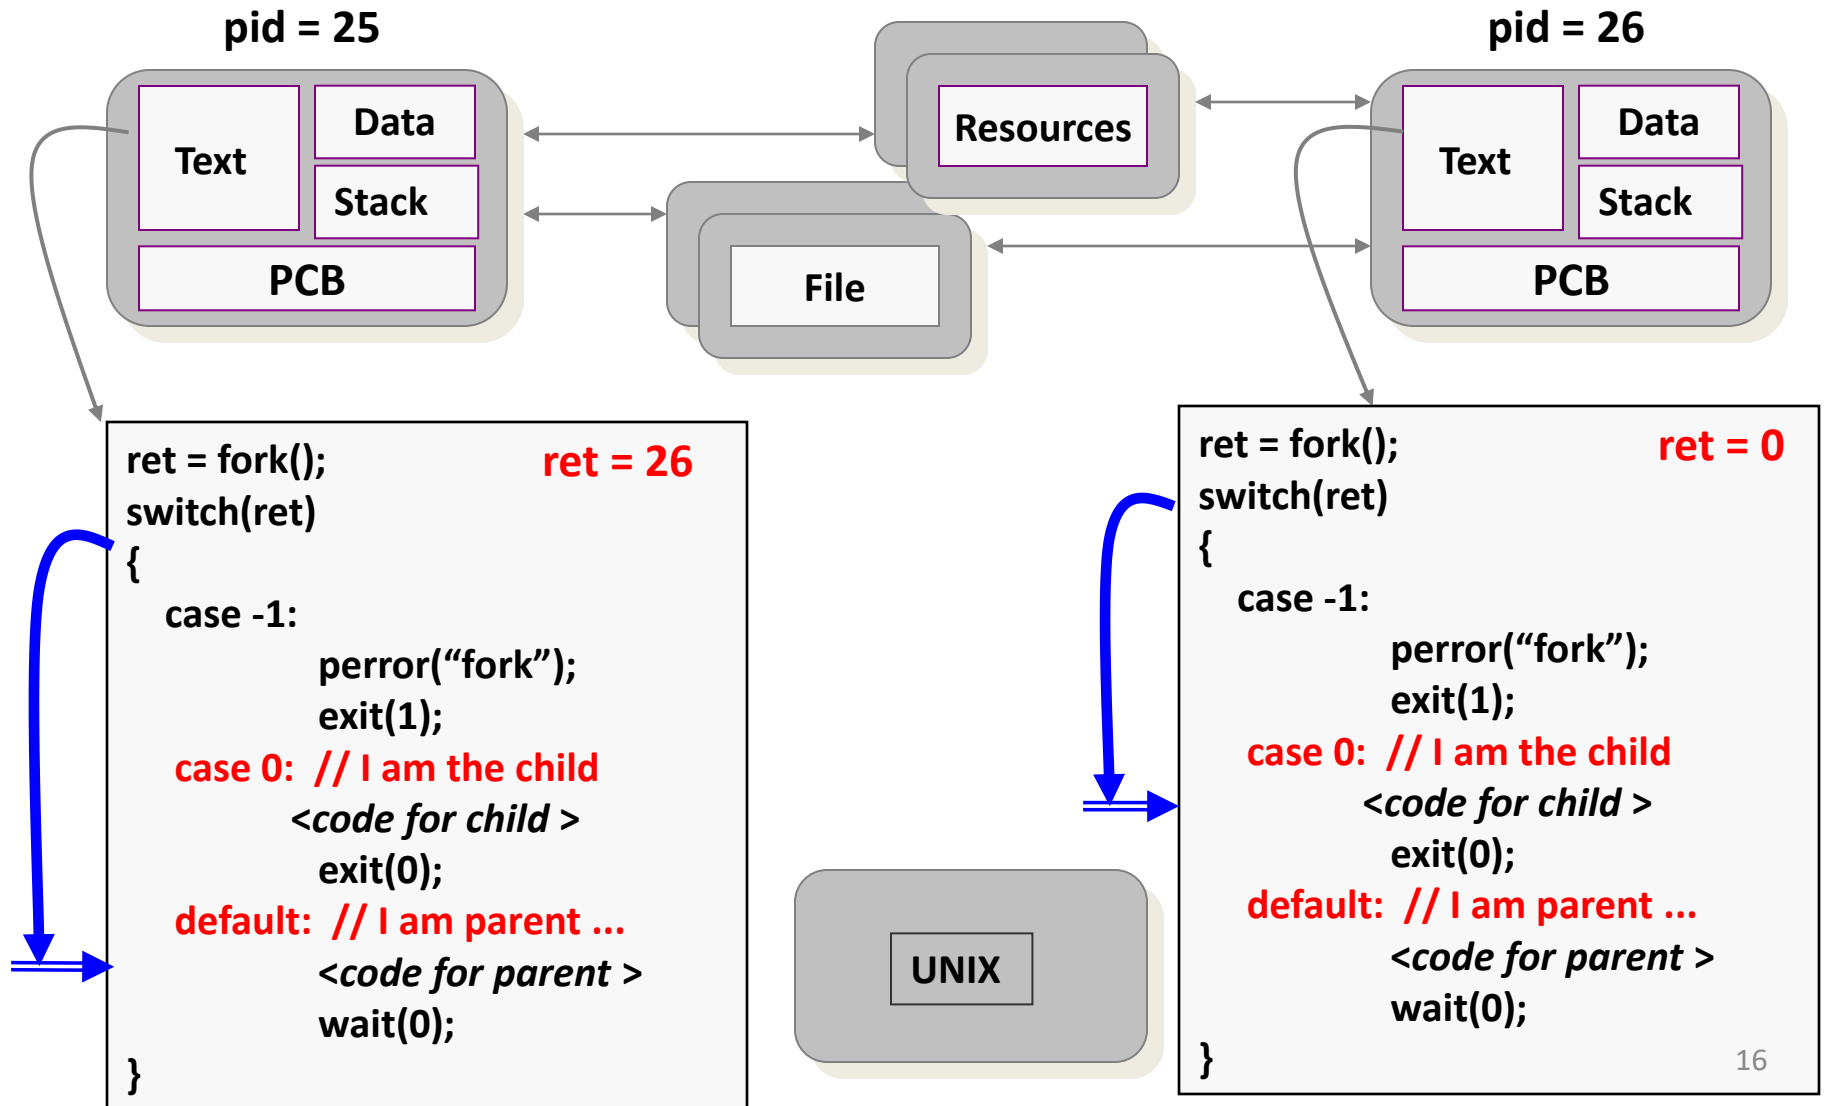
How fork Works (1)



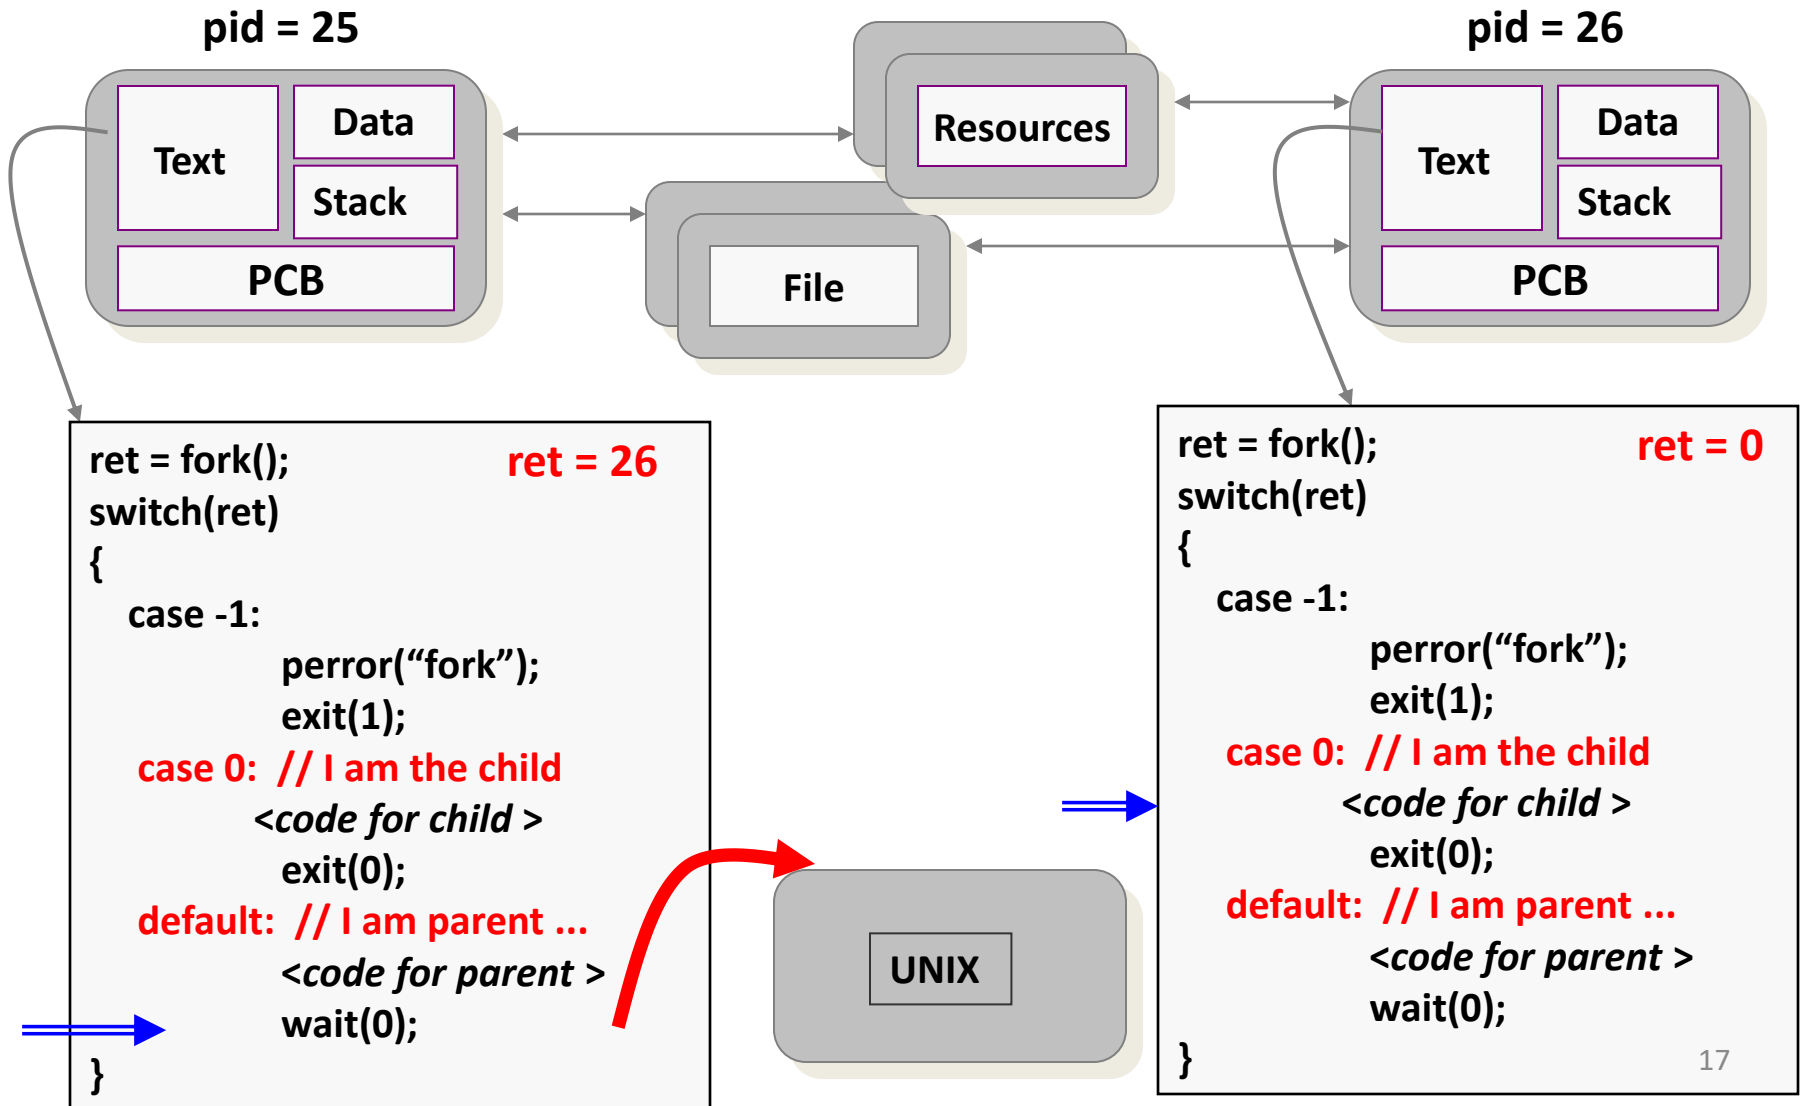
How fork Works (2)



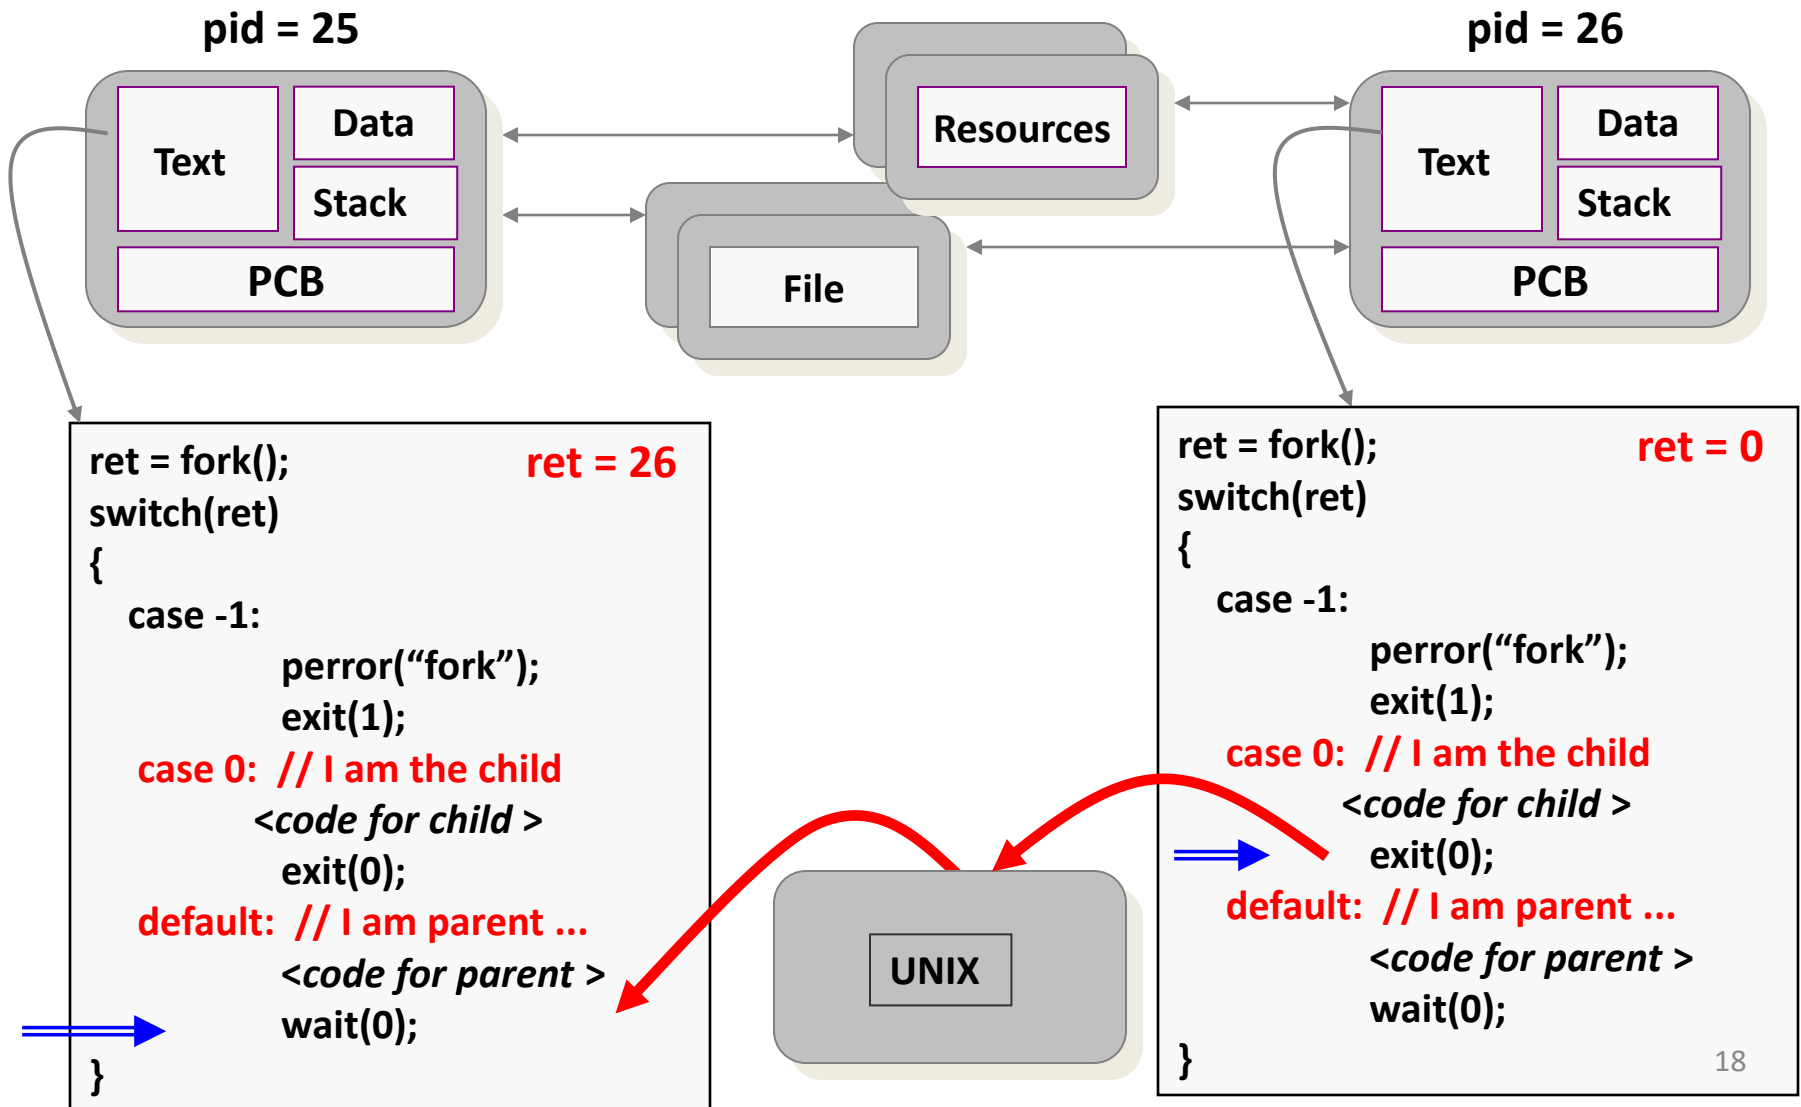
How fork Works (3)



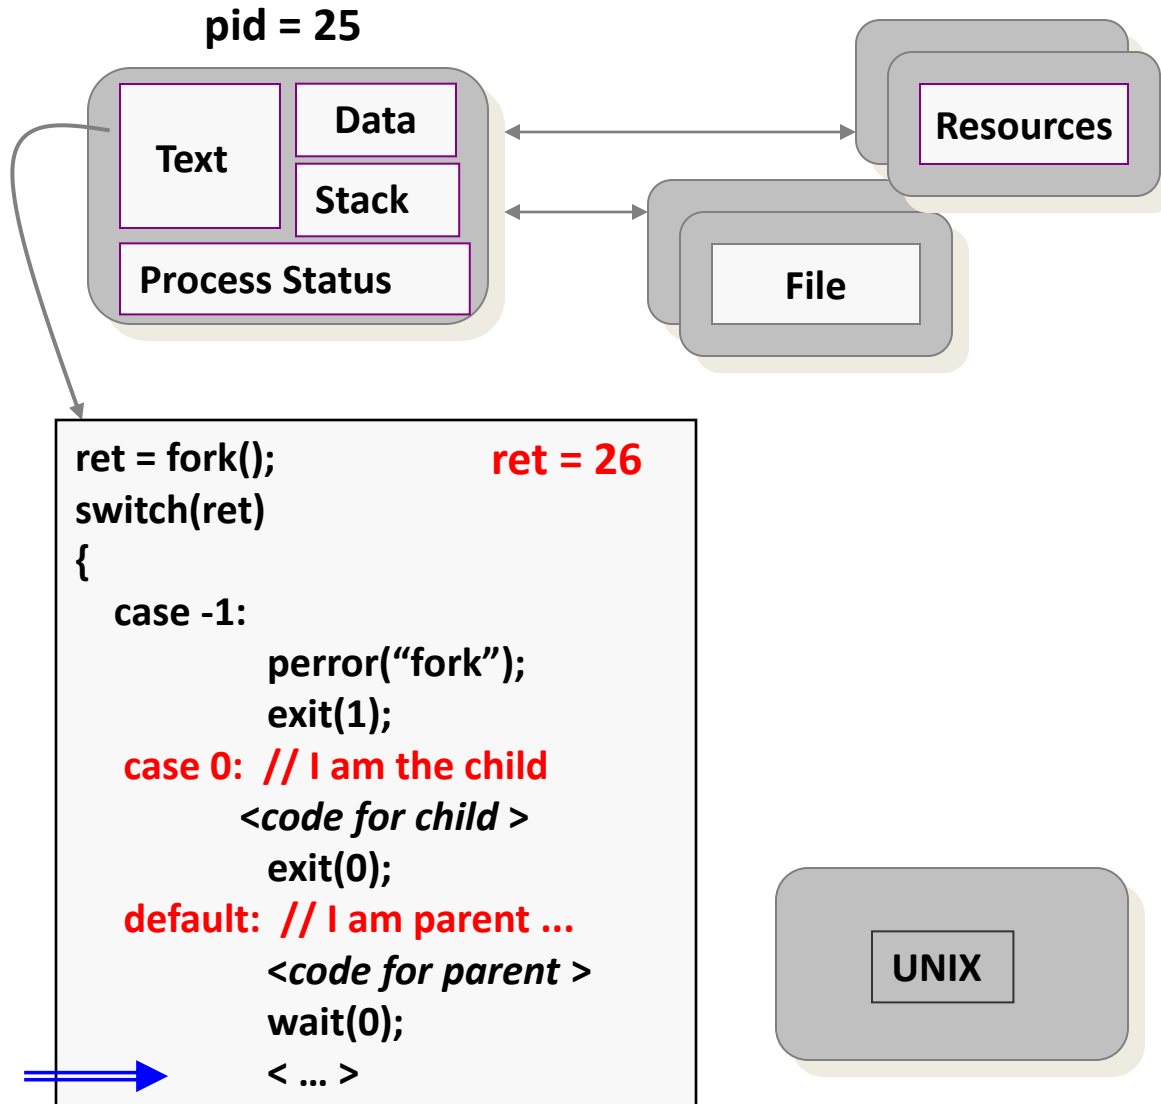
How fork Works (4)



How fork Works (5)



How fork Works (6)



Example: fork

```
#include <stdio.h>
```

```
int num = 0;
```

```
int main(){
```

```
    int pid;
```

```
    pid = fork();
```

parent: 0

```
    printf("%d", num);
```

```
    if(pid == 0) {          /*child*/
```

```
        num = 1;
```

```
    } else if(pid > 0) {    /*parent*/
```

```
        num = 2;
```

```
    }
```

```
    printf("%d", num);
```

```
}
```

child: 0

What if the child
calls fork() here?

parent: 2

child: 1

exit()

```
#include <stdlib.h>
```

```
void exit(int status);
```

man exit

- **status**: child's exit code (lower 8-bit only) returned to its parent
- Normal termination: closes all files, deallocates resources (memory, I/O buffer, etc.), removes child **zombie** procs, if any.
- Proc's PCB **may** or **may not** be released immediately
- If parent proc is blocked in `wait()`, unblock the parent proc & release child's PCB
- If parent proc hasn't called `wait()`, **holds the exit code in PCB until the parent calls `wait()`**
 - the child proc does not really die, but it enters a **zombie/defunct** state

wait()

```
#include <sys/types.h>
#include <sys/wait.h>
```

man wait

```
pid_t wait(int *status);
pid_t waitpid(pid_t pid, int *status, int options);
```

- Parent **may** want to wait for children to finish
 - Ex: a shell waiting for operations to complete
- **wait()**: wait for **any** child to terminate
 - Blocks until some child terminates
 - Returns the process ID of the child proc, and child's exit status
 - Or returns -1 if no children exist (i.e., already exited)
- **waitpid()**: wait for a **specific** child to terminate
 - Blocks till a child with particular proc ID terminates
- Both returns pid of the child proc that was terminated

What about
wait(0);

Example: exit, wait

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>

int main()
{
    pid_t fpid, pid, wpid;

    /* fork a child process */
    fpid = fork();
    if (fpid < 0) {
        /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }


    else if (fpid == 0) { /* child process */
        pid = getpid();
        printf("child: fpid = %d\n", fpid);
        printf("child: pid = %d\n", pid);
        exit(23);
    } else { /* parent process */
        int status = 0;
        pid = getpid();
        printf("parent: fpid = %d\n", fpid);
        printf("parent: pid = %d\n", pid);

        wpid = wait(&status);
        printf("parent: wait: wpid = %d, status = %d,
        exit code = %d\n", wpid, status,
        WEXITSTATUS(status));
    }
    return 0;
}
```

Process Termination

- **exit()**: normal proc termination
- **abort()**: abnormal termination
- **kill()**: terminate another proc
 - Permission: privileged user, or w/ the same user id
- **zombie process** (or defunct process): a proc has terminated, but its parent proc has not yet called **wait()**
 - PCB cannot be released until parent proc invokes **wait()**

man abort
man 2 kill
man 1 kill



How to remove
zombie process?

- **orphan process**: If a parent terminated w/o invoking **wait()**, child procs become orphan
 - OS assigns a **system** proc as the new parent
 - The new parent proc invokes **wait()** from time to time (**why?**)

Unix: init (pid 1)


Linux: init (pid may or may not be 1), upstart, etc.

exec()

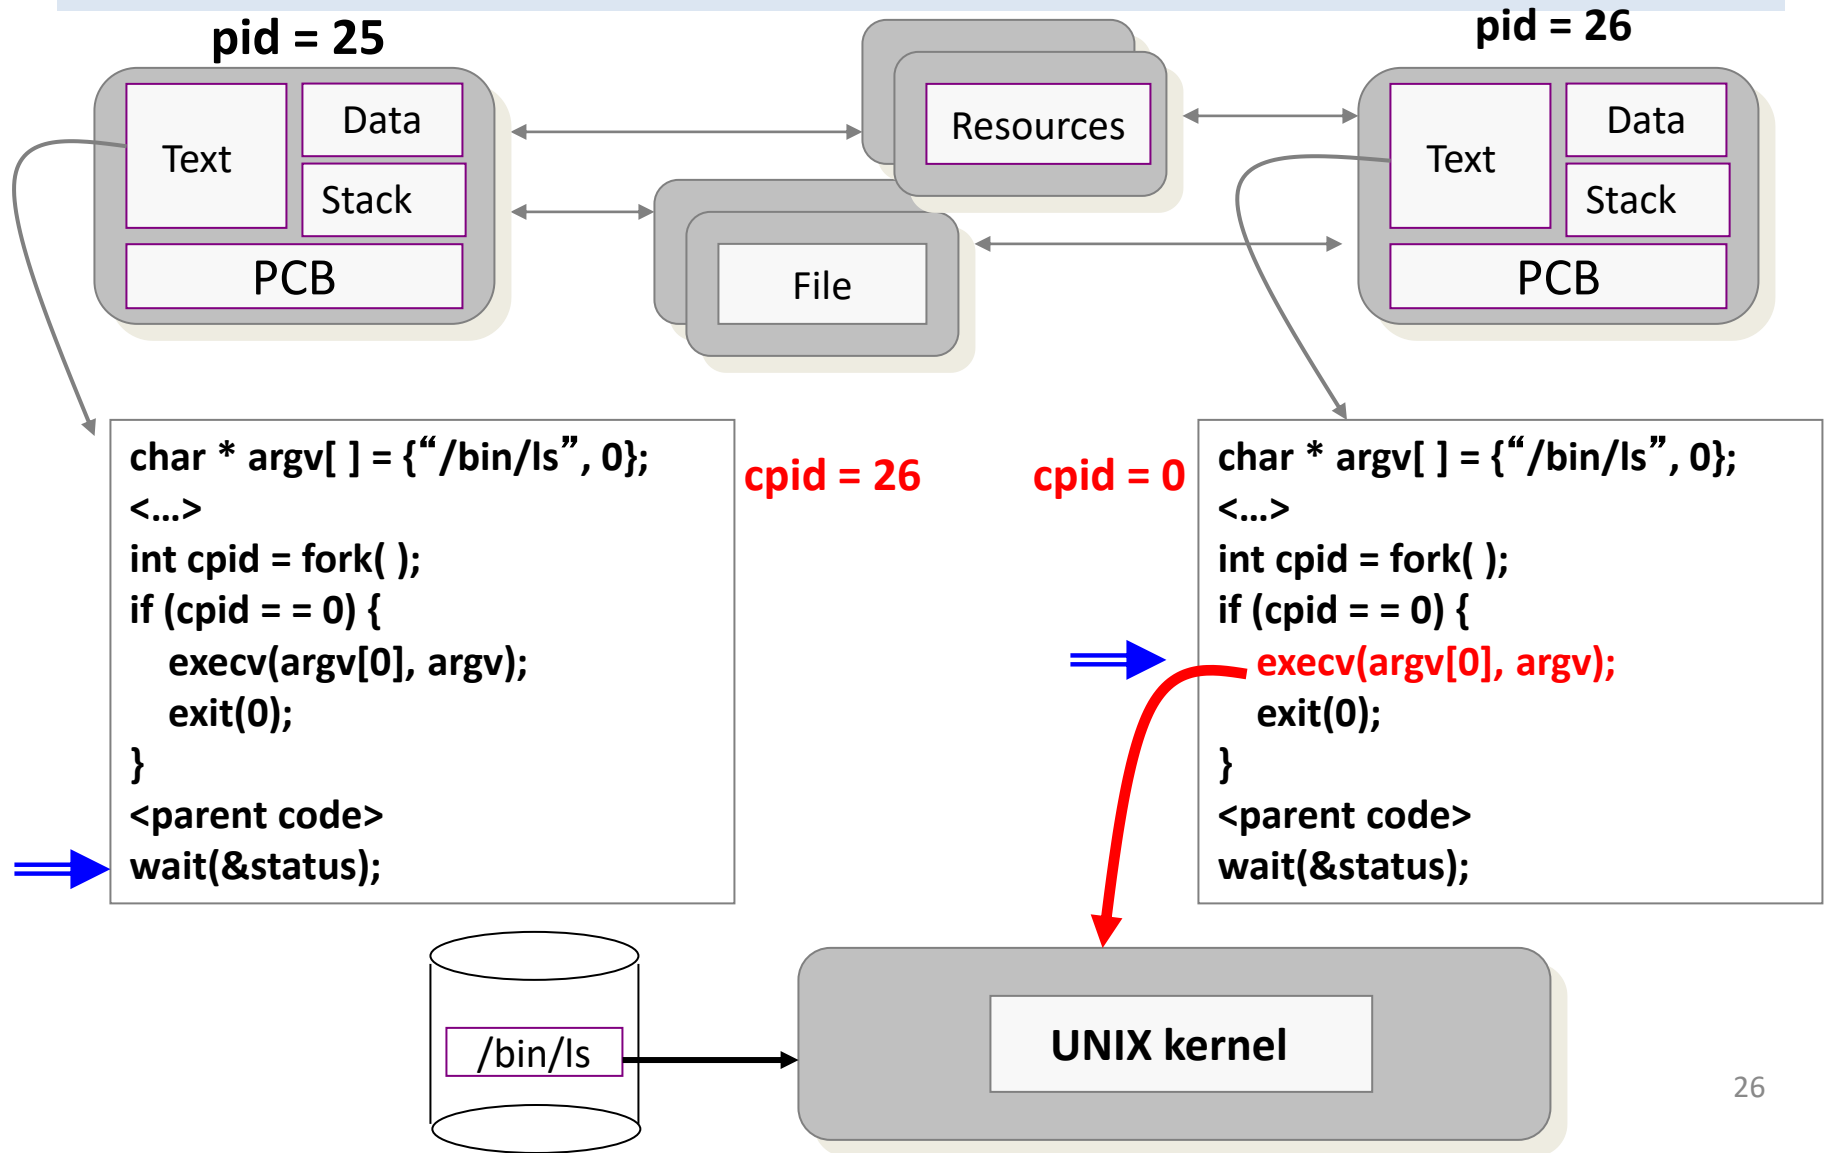
```
#include <unistd.h>
```

man 3 exec

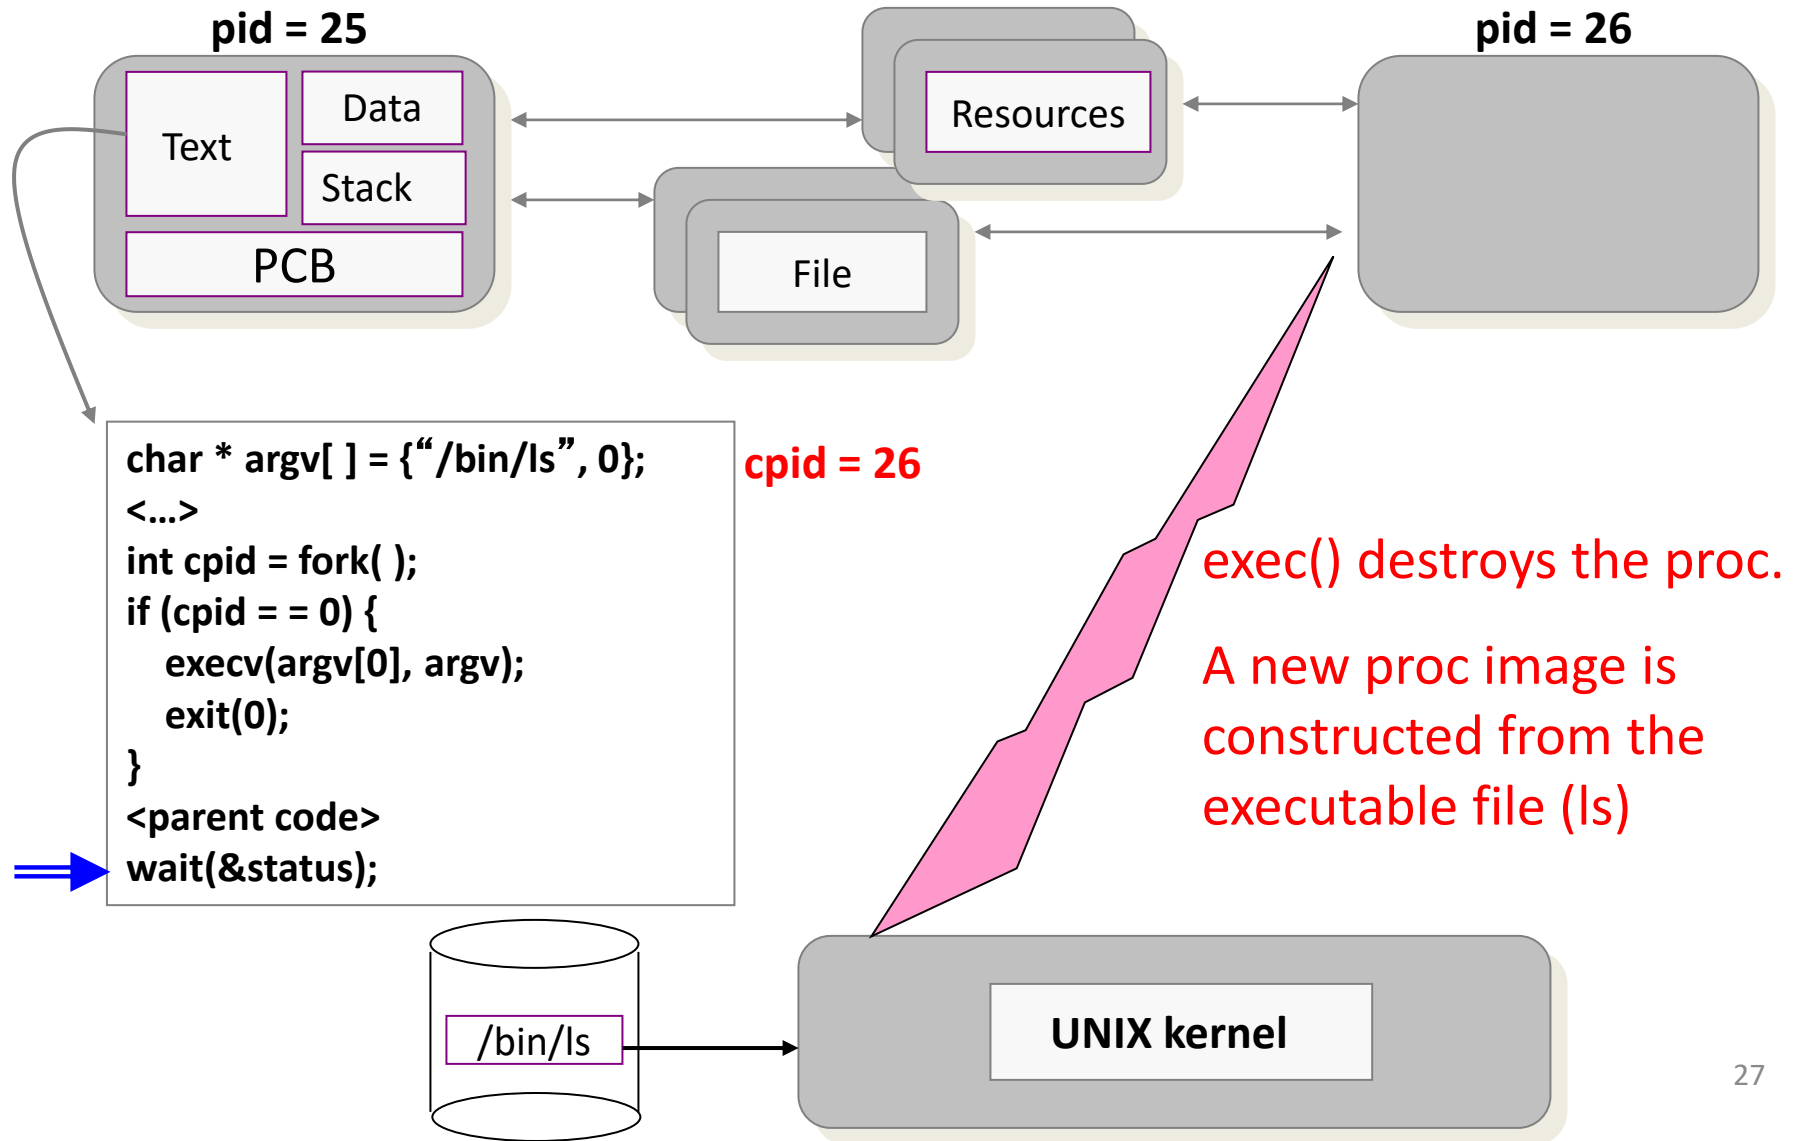
```
int execev(const char * path, char * const argv[]);
```

- executes a program - **replacing** the calling proc with a new proc
 - path: full path for the program to be executed
 - argv: the array of arguments for the program to execute
 - **each** argument is a **null-terminated** string
 - the first argument is the name of the program
 - the **last** entry in argv is **NULL** 
- After a successful exec, **no** return to the calling proc
 - calling proc replaced by the new proc
 - The new proc has the **same** pid (and parent pid) as the calling proc
- Return -1 only when error (calling proc still alive)

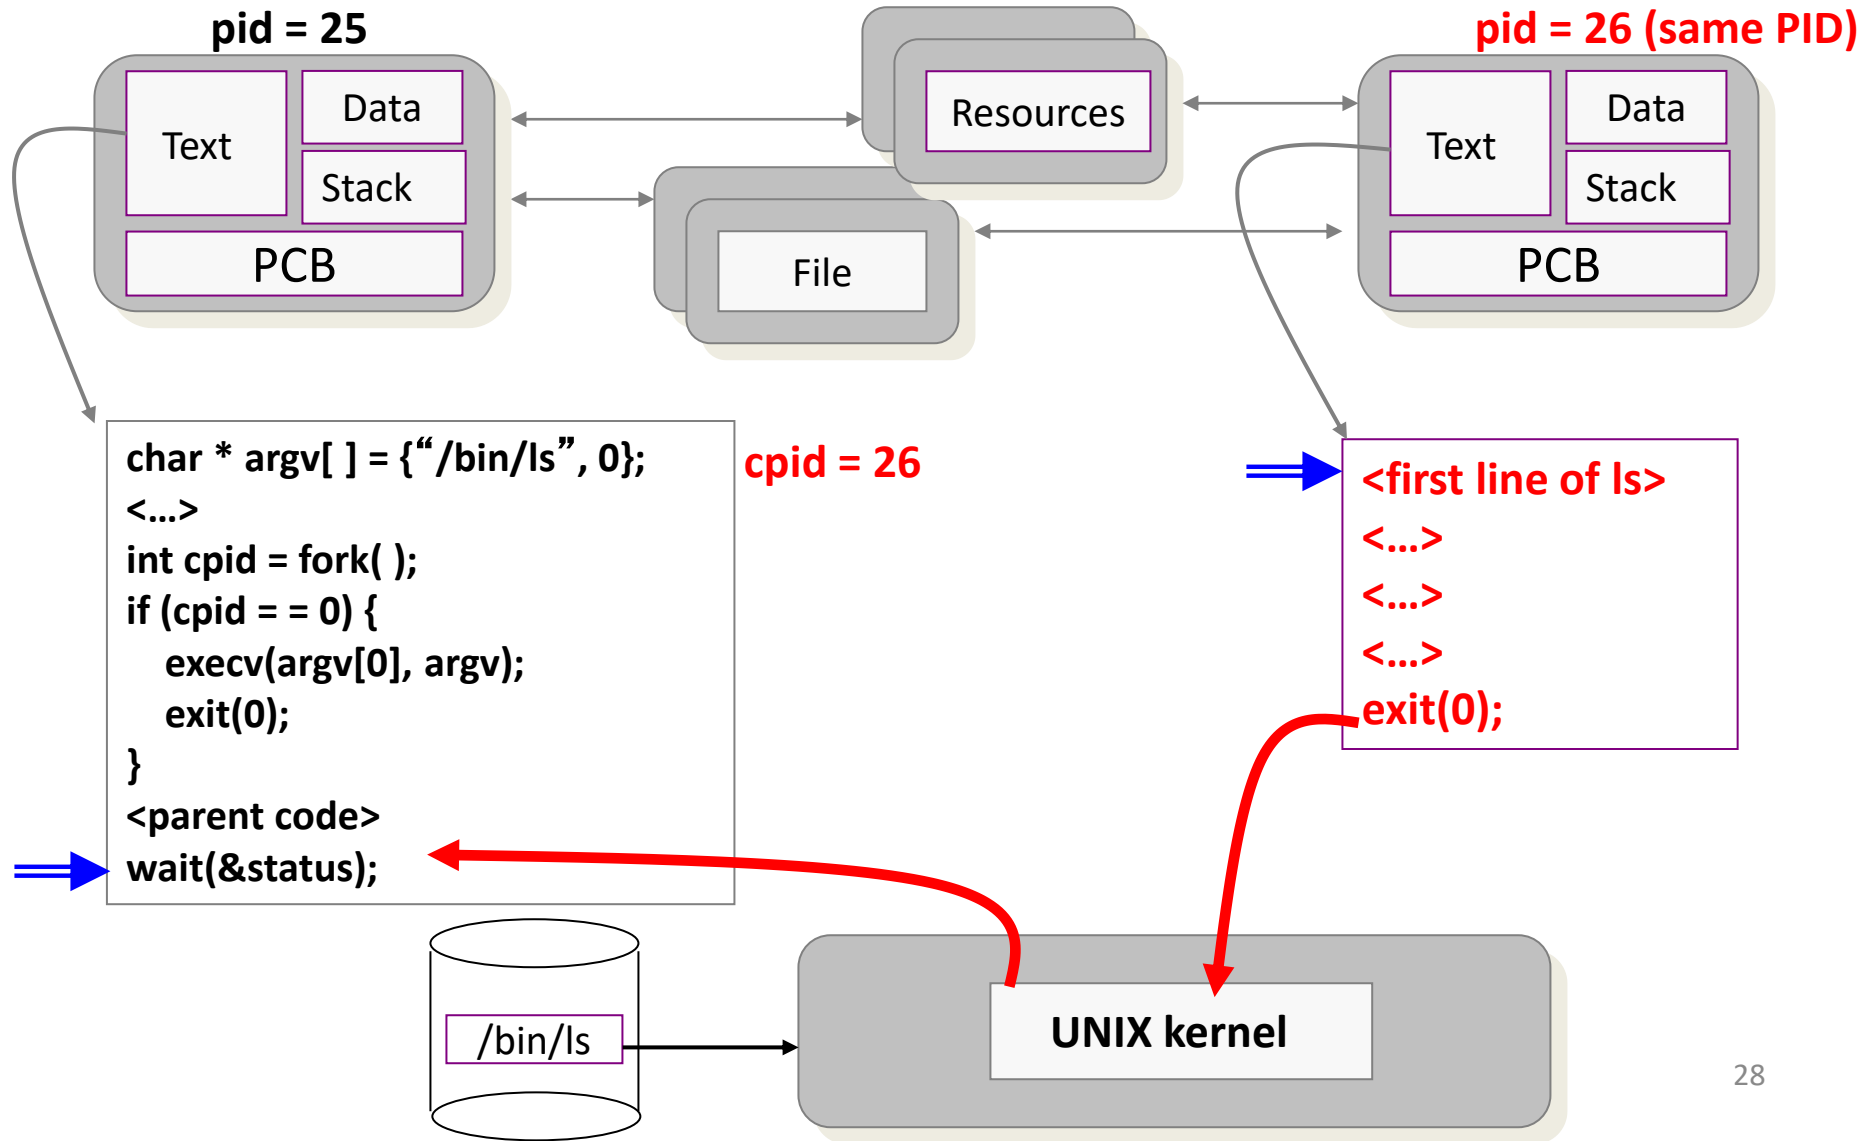
How execv Works (1)



How execv Works (2)



How execv Works (3)



execv Example

```
#include <stdio.h>
#include <unistd.h>

char * argv[] = {"/bin/ls", "-l", 0};
int main()
{
    int pid, status;

    if ( (pid = fork() ) < 0 ) {
        printf("Fork error \n");
        exit(1);
    }
    if(pid == 0) { /* Child executes here */
        execv(argv[0], argv);
        printf("Exec error \n");
        exit(1);
    } else { /* Parent executes here */
        wait(&status);
    }
    printf("Hello there! \n");
    return 0;
}
```

argv[0] = "/bin/ls";
argv[1] = "-l";
argv[2] = NULL;

Note the NULL string
at the end


Error handling

Error handling

Demo!

Example: A Simple Shell

- Shell is the parent process
 - E.g., `bash`
- Reads & parses cmd line
 - E.g., “`ls -l`”
- Invokes child proc
 - `fork`, `exec`
- Waits for child
 - `wait`
- Each Linux proc has three opened streams
 - `stdin`: standard input
 - keyboard input → `stdin`
 - `stdout`: standard output
 - `printf` → `stdout`
 - `stderr`: standard error





Fork Example 2 & 3: Spacetime Diagram

- Key Point: keep track of the execution of each process

```
void fork2()  
{  
    printf("L0\n");  
    fork();  
    printf("L1\n");  
    fork();  
    printf("Bye\n");  
}
```

```
void fork3()  
{ int i;  
  for (i=0; i<3; i++) {  
      fork();  
  }  
}
```



Fork Example 4 & 5: Spacetime Diagram

- How many processes? L0? L1? L2?
- The value of i in each process in spacetime diagram?

```
void fork4()
{   int i = 4;
    printf("L0\n");
    if (fork() != 0) {
        printf("L1\n");
        i = i + 2;
        if (fork() != 0) {
            printf("L2\n");
            fork();
        }
    }
    ++i;
    printf("Bye\n");
}
```

```
void fork5()
{   int i = 5;
    printf("L0\n");
    if (fork() == 0) {
        printf("L1\n");
        --i;
        if (fork() == 0) {
            printf("L2\n");
            fork();
            ++i;
        }
    }
    printf("Bye\n");
}
```

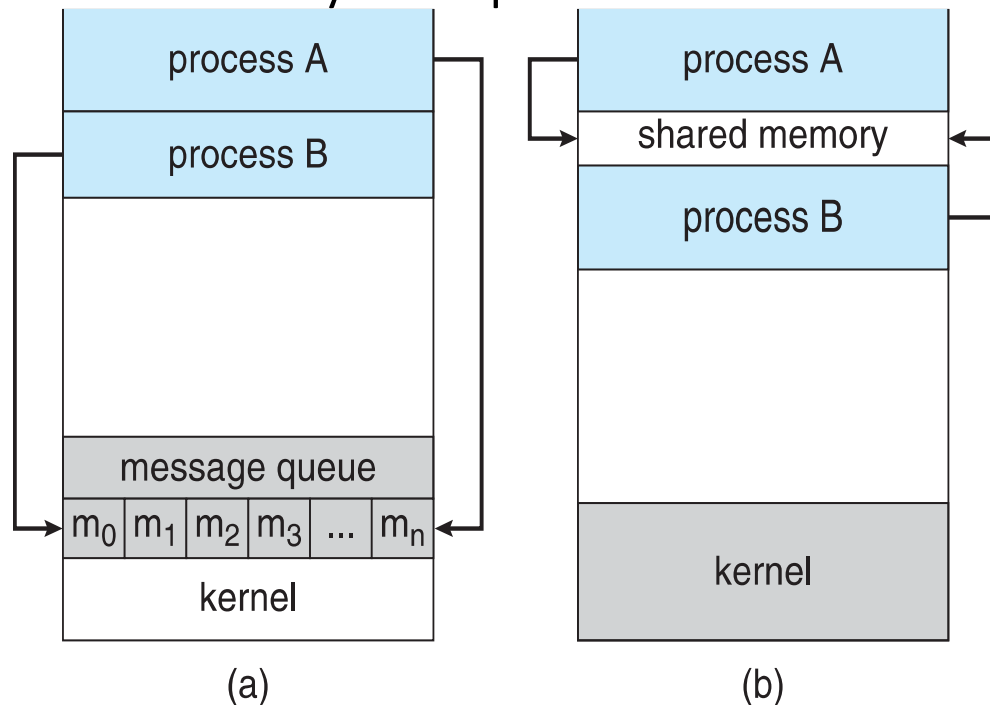

Multiprocess Architecture – Chrome Browser

- Many web browsers as single proc
 - If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser: multiproc w/ 3 types of procs:
 - Browser proc: manages UI, disk and network I/O
 - Renderer proc: 1 per website - render web pages HTML, Javascript
 - Plug-in proc: for each type of plug-in
 - Pros? Cons?



Interprocess Communication (IPC)

- Procs within a system may be independent or cooperating
 - **Independent** proc: cannot affect or be affected by other running proc
 - **Cooperating** proc: can affect or be affected by other procs
 - Info sharing
 - Computation speedup
- Cooperating procs need **IPC**
- Two models of IPC
 - **Message passing**: (a)
 - same or different computer
 - good for **smaller** amount of data exchange
 - **Shared memory**: (b)
 - same computer **only**
 - faster & more efficient for **large** data exchange
 - need **synchronization** & overhead (later)



IPC - Shared-Memory: bounded buffer

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;
item buffer[BUFFER_SIZE];
int in = 0; /* idx - producer */
int out = 0; /* idx - consumer */
```

- can only use BUFFER_SIZE-1 elements

Consumer

```
item next_consumed;
while (true) {
    while (in == out); /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next_consumed */
}
```

Producer

```
item next_produced;
while (true) {
    /* produce an item in next_produced */
    while (((in + 1) % BUFFER_SIZE) == out); /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```



Problems?

IPC - POSIX Shared Memory

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
```

Producer
gcc -lrt

```
int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr, "%s", message_0);
    ptr += strlen(message_0);
    sprintf(ptr, "%s", message_1);
    ptr += strlen(message_1);

    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
```

Consumer
gcc -lrt

```
int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

man shm_open
man mmap
man ftruncate

Demo!

parent/child
proc ?

Problems?

Sync & Async Message Passing

- Msg passing: a form of synchronization
- Msg passing: blocking or non-blocking
- **Blocking \equiv synchronous**
 - **Blocking send** -- the sender is blocked until msg is received
 - **Blocking receive** -- the receiver is blocked until msg is available
 - Involves proc state transition
- **Non-blocking \equiv asynchronous**
 - **Non-blocking send** -- the sender sends msg and continue
 - **Non-blocking receive** -- the receiver receives a valid msg or null msg
 - Why useful?
 - Do **not** have to involve proc state transition (why?)
- Different combinations possible
- I/O operations in OS: synchronous, asynchronous

Rendezvous


- Rendezvous: both send and receive are blocking
- Producer-consumer becomes trivial

Producer

```
message next_produced;  
while (true) {  
    /* produce an item in next_produced */  
    send(next_produced);  
}
```

Consumer

```
message next_consumed;  
while (true) {  
    receive(next_consumed);  
  
    /* consume the item in next_consumed */  
}
```



problems?
How to
resolve?


Link buffer

- Zero buffer (Rendezvous)
- Bounded buffer
- Unbounded buffer

Shared memory vs
msg passing?

IPC - Sockets

- Communication b/w a pair of sockets
- **socket**: port for communication – **local** to each host (IP)
 - IP:port or hostname:port
 - 161.25.19.8:1625: “port 1625 on host 161.25.19.8”
- Special IP 127.0.0.1 (**loopback**) - local system
- Well-known ports: < 1024 - smtp: 25, http: 80, https: 443, etc
- Socket types
 - **Connection-oriented (TCP)**
 - **Connectionless (UDP)**
 - MulticastSocket – multi recipients
- Linux cmd: netstat, lsof
- Binding: C, Java, etc.



Sockets in Java (Date server/client)

```
import java.net.*;
import java.io.*;
```

javac DateClient.java
java DateClient

```
public class DateClient
{
    public static void main(String[] args) {
        try {
            // could be changed one other than the localhost
            Socket sock = new Socket("127.0.0.1",6013);
            InputStream in = sock.getInputStream();
            BufferedReader bin =
                new BufferedReader(new InputStreamReader(in));

            String line;
            while( (line = bin.readLine()) != null)
                System.out.println(line);

            sock.close();
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

Demo!

```
import java.net.*;
import java.io.*;
```

javac DateServer.java
java DateServer

```
public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

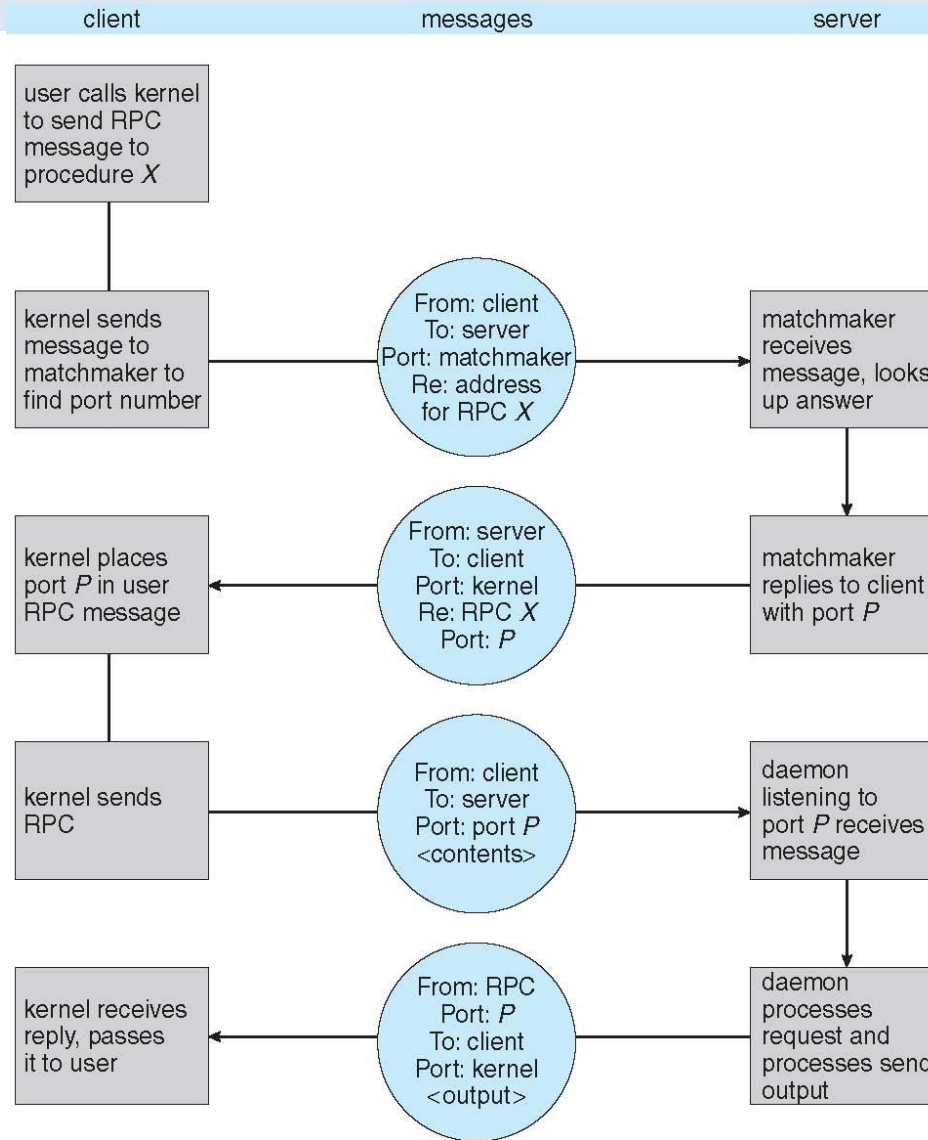
            /* now listen for connections */
            while (true) {
                Socket client = sock.accept(); ←
                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

problem?

IPC – Remote Procedure Call (RPC)



IPC - Ordinary Pipes

- Ordinary Pipes: communication in producer-consumer style
 - Unidirectional
 - Exists only when procs are communicating
 - Producer: writes to one end (the write-end of the pipe)
 - Consumer: reads from the other end (the read-end of the pipe)
 - Require parent-child relationship b/w communicating procs
 - Parent creates pipe (special file), then forks a child that shares pipe w/ parent
 - same machine only

- Shell: `ls | more`

How to program it in shell?

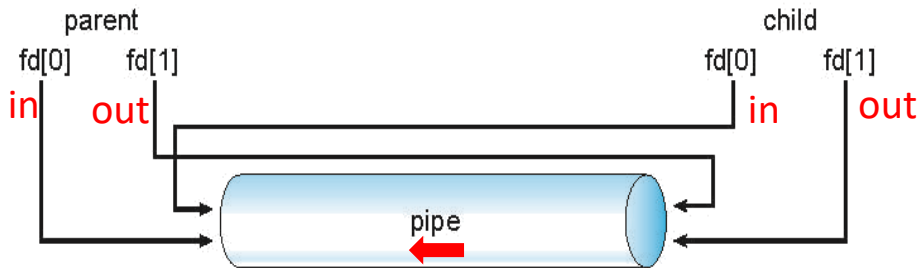
- stdout of `ls` becomes stdin of `more`
- Windows calls these anonymous pipes

Example: Ordinary Pipes

```
#include <sys/types.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
```

```
#define BUFFER_SIZE 25
#define READ_END 0
#define WRITE_END 1
```

```
int main(void)
{
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    int fd[2];
    pid_t pid;
```



```
/* create the pipe */
if (pipe(fd) == -1) {
    fprintf(stderr, "Pipe failed");
    return 1;
}

/* fork a child process */
pid = fork();

if (pid < 0) { /* error occurred */
    fprintf(stderr, "Fork Failed");
    return 1;
}

if (pid > 0) { /* parent process */
    /* close the unused end of the pipe */
    close(fd[READ_END]);

    /* write to the pipe */
    write(fd[WRITE_END], write_msg, strlen(write_msg)+1);

    /* close the write end of the pipe */
    close(fd[WRITE_END]);
}
else { /* child process */
    /* close the unused end of the pipe */
    close(fd[WRITE_END]);

    /* read from the pipe */
    read(fd[READ_END], read_msg, BUFFER_SIZE);
    printf("read %s", read_msg);

    /* close the write end of the pipe */
    close(fd[READ_END]);
}

return 0;
}
```

man pipe

IPC - Named Pipes

- Named Pipes
 - Bidirectional:
 - Half-duplex: $A \leftrightarrow B$; not simultaneously. E.g., walkie-talkie
 - Full-duplex: $A \leftrightarrow B$ simultaneously. E.g., telephone
 - No parent-child relationship is needed b/w communicating procs
 - >2 procs can use the named pipe for communication
 - Continue to exist even after procs have finished
- Named pipes in Linux/UNIX (aka FIFO)
 - As a file in FS; continue to exist until it is removed from FS
 - Bidirectional & half-duplex: for procs on the same machine
 - API: `mkfifo()`, `open()`, `read()`, `write()`, `close()`
- Named pipes in Windows
 - Bidirectional & full-duplex: for procs on local or remote machine
 - API: `CreateNamedPipe()`, `ConnectNamedPipe()`, `ReadFile()`, `WriteFile()`

Parallel Programming

- Get computations done in parallel
- Approaches
 - Single computer: multiple processes, multiple threads
 - Multiple computers: distributed system, cloud computing
- Issues
 - How many processes or threads?
 - Number of CPUs, nodes
 - How to partition the computation? Split data? data dependency?
 - Granularity of a process/thread
 - How to communication among processes/threads: IPC?
 - How to coordinate among processes/threads → synchronization (later)
 - Load balance?
 - Testing and debugging?

Question

- Q: [multi-choice] Which one is true?
 - A. Msg passing is better suited for large volume of data exchange than shared memory
 - B. IPC with shared memory is limited to parent-child processes only
 - C. Any two independent processes can communicate with ordinary pipe
 - D. RPC and socket allow communication across computers and within a computer
 - E. None of the above

Summary

- Process state transition diagram
- PCB content?
- What does context switch do?
- What is zombie process? Orphan process? How to resolve orphan process?
- Proc life cycle: create, terminate – fork(), exec(), exit(), wait()
- User enters “cc -o a b.c” in a shell, what does argv look like when shell invokes exec()?
- IPC
 - Shared memory
 - Message passing
 - RPC
 - Sockets
 - Pipes: ordinary pipe vs named pipe

Self Exercises

- 9/E: 3.1, 3.2, 3.4, 3.5, 3.9, 3.12, 3.13, 3.14, 3.15, 3.17, 3.18, 3.21, 3.25, 3.27, Project 1 part 1 & 2
- 10/E: 3.1, 3.2, 3.4, 3.5, 3.8, 3.10, 3.11, 3.12, 3.13, 3.14, 3.16, 3.17, 3.18, 3.19, 3.21, 3.26, Project 1 part 1~5